

THE REASONING CORPUS

Preserving Intellectual Provenance in the Age of AI



CRAIGS LLC
OCT 03, 2026

Author: Craig Medlin

Organization: CRAIGS LLC

Version: 1.1

Current Status: Observed implementation with a single user.

Validation: Microsoft Copilot + Work IQ + OneDrive.

Replication: Not yet independently reproduced.

Supporting development records, retrieval tests, and operational observations are preserved within the Reasoning Corpus. This whitepaper summarizes the resulting architecture and findings without requiring access to those internal records.

Executive Summary

Most digital systems preserve outcomes. They preserve emails, reports, presentations, photographs, documents, and final decisions.

What they often fail to preserve is the reasoning that produced those outcomes.

The Reasoning Corpus is a personal knowledge-preservation system designed to preserve observations, questions, interpretations, corrections, insights, and conclusions as they develop over time.

During development, an unexpected observation emerged. Procedures documented within the corpus were sometimes surfaced through repository retrieval before they were consistently reproduced through conversational continuity alone. This suggested that continuity may emerge from retrieval of preserved artifacts rather than memory alone.

The system combines:

- AI-assisted reasoning
- Voice-first journaling
- Structured metadata
- Markdown preservation
- Microsoft OneDrive
- Microsoft Copilot
- Work IQ grounding within Microsoft Copilot

The objective is not merely information retention.

The objective is continuity of understanding.

What did I know, when did I know it, why did I believe it, and how did that belief change?

1. The Problem

Human memory is incomplete.

Traditional archives suffer from a different limitation. They may preserve conclusions while losing:

- Context
- Reasoning
- Corrections
- Uncertainty
- Competing interpretations
- Intellectual development

Years later, a person may remember a conclusion but forget the reasoning that originally supported it.

This creates several challenges:

- Loss of rationale
- Hindsight distortion
- Broken continuity
- Difficulty reconstructing decisions
- Difficulty tracing idea evolution

The corpus was created in response to that problem. A foundational observation emerged early:

Preserved context $\neq$ Retrieved context
Information exists $\neq$ Information is visible

Storage alone does not guarantee retrieval. Retrieval alone does not guarantee understanding.

2. The Hypothesis

Reasoning can be archived.

If conversations, journals, corrections, reflections, and analyses are preserved with sufficient fidelity and metadata, future retrieval systems may reconstruct not only facts, but intellectual evolution.

Instead of preserving only:

Conclusion

the corpus attempts to preserve:

```
graph TD; A[Observation] --> B[Interpretation]; B --> C[Challenge]; C --> D[Correction]; D --> E[Refinement]; E --> F[Conclusion];
```

Observation
↓
Interpretation
↓
Challenge
↓
Correction
↓
Refinement
↓
Conclusion

The path may ultimately be more valuable than the destination.

3. Current Tested Environment

The Reasoning Corpus has been developed and observed in the following environment:

- Microsoft Copilot

- Microsoft 365
- Microsoft OneDrive
- Work IQ grounding within Microsoft Copilot
- Markdown files
- Structured metadata
- Direct-download ingestion
- Voice-first journaling

Work IQ was active in the Microsoft Copilot environment in which the corpus was developed and tested. This paper does not evaluate Work IQ independently or claim that Work IQ alone produced the observed retrieval results.

3.1 Validation Boundary

The corpus has not been systematically tested using:

- ChatGPT
- Claude
- Gemini
- Local language models
- Alternative cloud repositories
- Alternative retrieval systems
- Microsoft Copilot with Work IQ unavailable or inactive

Therefore, this paper does not claim universal portability, identical behavior across AI models, or equivalent retrieval quality outside the Microsoft ecosystem.

The Reasoning Corpus has demonstrated practical usefulness within the Microsoft Copilot, Work IQ, OneDrive, and Markdown environment in which it was developed.

4. Architectural Model

The corpus consists of six layers.

Layer 1: Primary Reasoning

Human-AI conversations provide raw reasoning material. Examples include:

- Questions
- Reactions
- Interpretations
- Corrections
- Experiments
- Emerging frameworks
- New observations

The purpose of conversation is discovery. Conversation itself is not treated as the authoritative artifact.

Layer 2: Structured Handoffs

Significant conversations are converted into structured Markdown artifacts.

Filename convention:

YYYY-MM-DD_HHMM_Topic_Name.md

Example:

2026-09-30_1238_Corpus_Risk_And_Long_Term_Vision.md

The purpose of the handoff is not transcript preservation. The purpose is reasoning preservation.

Layer 3: Voice-First Journaling

The corpus includes personal journal entries created outside AI conversations.

Workflow:

```
graph TD; A[Spoken reflection] --> B[Mobile voice recording]; B --> C[.m4a export]; C --> D[.txt export]; D --> E[Markdown conversion]; E --> F[Corpus artifact];
```

Spoken reflection
↓
Mobile voice recording
↓
.m4a export
↓
.txt export
↓
Markdown conversion
↓
Corpus artifact

The recording application exports:

- Audio (.m4a)
- Transcript (.txt)

The transcript is converted into a structured Markdown artifact. Three preservation layers are maintained:

```
graph TD; A[Audio] --> B[Transcript]; B --> C[Structured Markdown record];
```

Audio
↓
Transcript
↓
Structured Markdown record

Each layer serves a different purpose. The original audio is the closest preserved form of the spoken entry. The transcript provides searchable text. The Markdown version adds structure, meaning, and retrieval cues.

Layer 4: Canonical Preservation Format

Markdown (.md) is the canonical textual preservation format.

Markdown was selected because it is:

- Human-readable
- Machine-readable
- Searchable
- Portable
- Durable
- Easy to migrate
- Independent of proprietary software

Markdown is not merely a file format. It is a long-term preservation strategy.

Layer 5: Repository Preservation

The corpus is stored within a structured OneDrive repository.

Current architecture:

CHAT_HANDOFF_CORPUS

- |— Archive
- |— CANON
- |— Derivatives
- |— Indexes
- |— Observations
- |— Reference
- |— Validation

Within CANON:

CANON

- |— Essays
- |— Journal
- |— Milestones
- |— Notes
- |— Records

Within Records:

Records

- |— YYYY
 - |— Active
 - |— Foundational

The structure supports preservation, classification, validation, and retrieval.

Layer 6: Retrieval

The retrieval layer allows preserved records to participate in future reasoning. The objective is not merely finding files. The objective is reconstructing understanding.

Typical retrieval questions include:

- When did this idea originate?
 - What evidence supported it?
 - When did I change my position?
 - Which discussions influenced this framework?
 - What objections were considered?
-

5. Canonical Artifact Formats

5.1 Chat Handoff

Chat Handoff

Date:

Time:

Primary Topic

Why This Conversation Mattered

Why This Matters

Key Insights

People

Retrieval Keywords

Key Outcomes

Continuity Links

The Reasoning Corpus does not currently rely on formal graph relationships. Continuity emerges through overlapping retrieval signals including timestamps, retrieval keywords, metadata sections, continuity links, and repeated references. Together these elements create multiple retrieval paths to the same concepts.




Figure 1. Overlapping semantic cues create a graph-like retrieval structure without requiring explicit graph construction.

Why Two Significance Sections Exist

“Why This Conversation Mattered” explains why the specific discussion was preserved.

“Why This Matters” explains the broader lesson or implication.

The first preserves the event. The second preserves the principle.

Timestamp Authority and Date-Boundary Integrity

The Date and Time fields describe when the handoff artifact was created. They do not necessarily describe when the conversation began or when the events summarized within it occurred.

This distinction becomes important when a conversation spans multiple calendar dates. A discussion may begin on one date, continue beyond midnight, and produce a handoff on the following date.

Conversation Date
≠
Artifact Creation Date

For handoff metadata, the artifact creation timestamp is authoritative.

The filename, Date field, and Time field must describe the same creation event:

Filename:
2026-10-02_2347_Date_Boundary_Integrity.md

Date:
2026-10-02

Time:
23:47 EDT

Earlier conversation dates may be preserved within the summary when relevant. They must not replace the artifact creation date in the filename or metadata.

This produces a broader provenance rule:

Provenance fields describe the artifact, not the topic.

Before a handoff is considered complete, the following values should be checked for agreement:

- Local artifact creation date
- Artifact creation time
- Filename timestamp
- Internal Date field
- Internal Time field and stated time zone

This verification is particularly important when a conversation spans midnight or when a handoff is generated late in the evening.

5.2 Journal Entry

Journal Entry

Date:

Time:

Source:

Original Audio:

Original Transcript:

Primary Topic

Summary

Main Reflections

Key Observations

Notable Statements

Retrieval Keywords

Related Topics

Why This Entry Mattered

Why This Matters

Chat handoffs preserve reasoning developed through dialogue. Journal entries preserve reasoning captured through personal reflection.

6. Direct-Ingestion Workflow

One of the simplest improvements was changing the browser's default download location to a folder inside the synchronized corpus.

Original workflow:

```
graph TD; A[Copilot] --> B[Download]; B --> C[Downloads folder]; C --> D[Manual move]; D --> E[Corpus]
```

Optimized workflow:

```
graph TD; A[Copilot] --> B[Download]; B --> C[Corpus]; C --> D[OneDrive sync]
```

↓
Retrieval surface

Benefits include:

- Reduced friction
 - Fewer manual steps
 - Faster synchronization
 - Lower probability of misplaced records
 - More consistent archival behavior
-

7. Core Design Principles

Record Fidelity

Record fidelity is more important than organization. An imperfect filename may be inconvenient. An incomplete record may be misleading.

The greatest risk is degradation of:

- Context
- Chronology
- Corrections
- Provenance
- Reasoning pathways

Intellectual Provenance

The corpus preserves:

- Evidence
- Questions
- Challenges
- Corrections
- Alternatives considered
- Context
- Evolution of understanding

The objective is to preserve how conclusions emerged.

Human-Readable Authority

AI retrieval is useful, but the AI is not the authority over the record.

For chat handoffs, the preserved Markdown record is the canonical textual artifact. For journal entries, the Markdown record is the canonical retrieval artifact, while the original audio and transcript remain supporting source records.

Lightweight Metadata

The corpus depends on a small set of stable metadata fields:

- Topic
- Significance
- Insights
- Keywords
- Outcomes
- Continuity

The metadata exists to preserve meaning, not bureaucracy.

Chronological Integrity

A reasoning-provenance system must distinguish between the chronology of the subject and the chronology of the record.

A handoff may describe events that occurred earlier while remaining a new artifact created at a later time. Both timelines may be valid, but they answer different questions:

Conversation chronology:

When did the discussion or event occur?

Artifact chronology:

When was the preserved record created?

Confusing these timelines weakens provenance because the record no longer clearly identifies its own creation event.

The corpus therefore treats the artifact creation timestamp as authoritative for filenames and metadata. Relevant conversation dates remain part of the narrative record.

Chronological integrity does not require every date in an artifact to be identical. It requires each date to have a clear and accurate meaning.

Retrieval Over Recall

The corpus is built on a practical observation:

▮ *Reliable retrieval may be more valuable than perfect memory.*

8. Observed Challenges

Downloadability Drift

A recurring issue involved generation of handoff content without generation of a downloadable Markdown artifact.

Repeated correction was required:

■ *"Downloadable please."*

The lesson:

Correct content ≠ Completed artifact

Naming Drift

Naming conventions evolved over time. Although retrieval continued to function, inconsistency created:

- Visual noise
- Additional review effort
- Reduced predictability

This reinforced the value of a standardized naming convention.

Metadata Drift

Required sections were occasionally omitted or applied inconsistently. The distinction between Why This Conversation Mattered and Why This Matters was one example.

Date-Boundary Drift

A date-boundary failure emerged during a late-evening handoff workflow.

A conversation had begun on an earlier date and contained numerous references associated with that date. After the calendar date changed, a new handoff inherited the dominant date from the conversation rather than the date on which the artifact was created.

This exposed two simultaneous timelines:

Conversation context date
≠
Artifact creation date

Neither timeline was inherently incorrect. The failure occurred because the conversation date was used where the artifact creation date was required.

The resulting handoff contained internally plausible metadata that did not accurately describe its own creation. Manual review identified the discrepancy, and the artifact was corrected.

The incident produced several operational requirements:

- Handoff metadata must reflect artifact creation.
- The filename date and internal Date field must match.
- The Time field must represent the same creation event.
- Earlier conversation dates should be preserved in the summary when relevant.
- Handoffs created near midnight require an explicit local-date check.
- A handoff is not complete until its timestamp fields have been verified.

The broader lesson is:

▮ Provenance fields describe the artifact, not the topic.

Date-Boundary Drift is related to Metadata Drift, but it represents a distinct failure mode. Metadata Drift concerns missing or inconsistently applied fields. Date-Boundary Drift occurs when two valid timelines exist and the wrong timeline is assigned authority over the artifact.

The correction followed the corpus's established development pattern:

```
graph TD; A[Observation] --> B[Error]; B --> C[Correction]; C --> D[Standard];
```

Observation
↓
Error
↓
Correction
↓
Standard

Instruction Retention Variability

Procedures agreed upon in prior conversations were not always applied consistently. Examples included:

- Missing downloads
- Missing metadata sections
- Naming inconsistencies

This produced an important lesson:

▮ *Conversational agreement is not a reliable process-control mechanism.*

Standards should exist inside retrievable records.

Indexing Delay

The system must distinguish:

File created ≠ File synced ≠ File indexed ≠ File retrieved

A file may exist before it becomes discoverable.

Survivor Bias

The corpus preserves what was recorded. It does not preserve everything that occurred. Future readers must not confuse archived history with complete history.

9. The OneDrive Context Revelation

One of the most interesting findings emerged when repository retrieval surfaced context before conversational continuity reproduced it reliably.

The informal observation became:

▮ *“OneDrive remembered first.”*

This should not be interpreted literally. The evidence supports a narrower conclusion:

▮ *Repository retrieval sometimes recovered documented procedures before conversational continuity consistently reproduced them.*

Conceptually:

Conversational continuity

=

What remains available in the conversation

Repository retrieval

=

What can be rediscovered from preserved artifacts

This may be one of the most important architectural findings of the corpus. The archive became part of the cognitive system.

Craig $\rightleftharpoons$ AI

evolved into

Craig $\rightleftharpoons$ AI $\rightleftharpoons$ Corpus

The corpus stopped functioning solely as storage. It began functioning as continuity.

10. Long-Term Vision

The Reasoning Corpus is not intended to become a database of answers. It is intended to become a historical record of thought.

If maintained for decades, it may preserve:

- Origins of ideas
- Changes in belief
- Failures
- Corrections
- Framework development
- Personal reflections
- Intellectual milestones

- Evolution of language
- Evolution of understanding

Most archives preserve conclusions. The Reasoning Corpus attempts to preserve the path.

Its value may not be helping someone remember what they knew. Its value may be helping someone understand how they came to know it.

Key Findings

1. Reasoning can be preserved as a retrievable artifact.
 2. Retrieval quality depends heavily on record fidelity.
 3. Markdown provides a practical long-term preservation format.
 4. Repository retrieval sometimes restored context before conversational continuity reproduced it.
 5. Continuity may emerge from preserved artifacts rather than memory alone.
 6. Accurate provenance requires distinguishing between event chronology, conversation chronology, and artifact chronology.
-

Conclusion

The Reasoning Corpus is an attempt to preserve a layer that is often lost: reasoning itself.

The implementation is technically simple:

- Voice recordings
- Transcripts
- Markdown artifacts
- Structured metadata
- OneDrive storage
- Microsoft Copilot retrieval
- Work IQ-assisted context discovery

The complexity lies not in technology. The complexity lies in maintaining continuity.

The central finding to date is:

Reliable AI continuity does not necessarily reside inside the AI. It can emerge from faithful records that remain available for retrieval.

The corpus therefore functions as more than an archive. It functions as a reasoning-provenance system.

The development of the corpus also revealed that continuity depends on preserving distinctions between different kinds of chronology. An event, a conversation about the event, and the artifact produced from that conversation may each have a different timestamp. Accurate provenance requires those timestamps to remain distinguishable rather than collapsing them into a single date.

The corpus therefore preserves not only what a record describes, but also when the record itself became part of the historical record.

This distinction reinforces the corpus's recursive character. Its records preserve the reasoning, corrections, standards, and publication history involved in developing the system itself.

| The corpus itself contains the history of creating the corpus.